

Meeting Report

September 25, 2019

This report is mainly talking about the academic communication with Prof. YanLin Geng at Sunday, September 22, 2019.

1 Some variations of GANs

In this section we will introduce some variations of GANs.

1.1 InfoGAN

The InfoGAN [1] bind information theory to GANs.

Abstract

This paper describes InfoGAN, an information-theoretic extension to the Generative Adversarial Network that is able to learn disentangled representations in a completely unsupervised manner. InfoGAN is a generative adversarial network that also maximizes the mutual information between a small subset of the latent variables and the observation. We derive a lower bound of the mutual information objective that can be optimized efficiently. Specifically, InfoGAN successfully disentangles writing styles from digit shapes on the MNIST dataset, pose from lighting of 3D rendered images, and background digits from the central digit on the SVHN dataset. It also discovers visual concepts that include hair styles, presence/absence of eyeglasses, and emotions on the CelebA face dataset. Experiments show that InfoGAN learns interpretable representations that are competitive with representations learned by existing supervised methods.

Problems The vanilla GAN imposes no restriction on input noise z , thus learns highly entangled representations.

Methods Decompose the input noise into two parts: 1) incompressible noise; 2) latent code.

Contributions Provides more interpretability by introducing latent codes on input noise; learns disentangled representations in a completely unsupervised manner.

1.2 SS-InfoGAN

The ss-InfoGAN [2] use a relatively small amount of labels to guide the training of InfoGAN.

Abstract

In this paper we propose a new semi-supervised GAN architecture (ss-InfoGAN) for image synthesis that leverages information from few labels (as little as 0.2%, max. 10% of the dataset) to learn semantically meaningful and controllable data representations where latent variables correspond to label categories. The architecture builds on Information Maximizing Generative Adversarial Networks (InfoGAN) and is shown to learn both continuous and categorical codes and achieves higher quality of synthetic samples compared to fully unsupervised settings. Furthermore, we show that using small amounts of labeled data speeds-up training convergence. The architecture maintains the ability to disentangle latent variables for which no labels are available. Finally, we contribute an information-theoretic reasoning on how introducing semi-supervision increases mutual information between synthetic and real data.

Problems InfoGAN [1] have recently been shown to learn disentangled representations, yet the extracted representations are not always directly interpretable by humans and lack direct measures of control due to the unsupervised training scheme.

Methods Leverage information from few labels, by decomposing latent code into two parts: unsupervised and supervised. Direct training to increase $I(C_{ss}; X)$ and $I(C_{ss}; \tilde{X})$, result in the increase of $I(X; \tilde{X})$.

Contributions Achieve higher quality of synthetic samples compared to fully unsupervised settings; faster convergence rate than InfoGAN.

1.3 SemiGAN

The SemiGAN [3] is an extension of semi-supervised learning of GAN.

Abstract

We extend Generative Adversarial Networks (GANs) to the semi-supervised context by forcing the discriminator network to output class labels. We train a generative model G and a discriminator D on a dataset with inputs belonging to one of N classes. At training time, D is made to predict which of $N + 1$ classes the input belongs to, where an extra class is added to correspond to the outputs of G . We show that this method can be used to create a more data-efficient classifier and that it allows for generating higher quality samples than a regular GAN.

Problems The vanilla GAN is suitable for classification, and how the generator G , discriminator D and classifier C will affect each other.

Methods Use softmax instead of sigmoid.

Innovations

- Use GANs in semi-supervised classification.
- Better performance than CNN on some datasets.
- Faster training and better outputs of generator.

Future work Use ladder network for D and G .

1.4 CGAN

The CGAN [4] proposes a conditional version of GAN.

Abstract

Generative Adversarial Nets were recently introduced as a novel way to train generative models. In this work we introduce the conditional version of generative adversarial nets, which can be constructed by simply feeding the data, y , we wish to condition on to both the generator and discriminator. We show that this model can generate MNIST digits conditioned on class labels. We also illustrate how this model could be used to learn a multi-modal model, and provide preliminary examples of an application to image tagging in which we demonstrate how this approach can generate descriptive tags which are not part of training labels.

Problems Vanilla GAN can not generate a specific class of samples we want; vanilla GAN is not scalable for big amount of data categories; plenty of works focus on one-to-one mapping, however, one-to-many (e.g., image labeling) is also a common problem.

Methods Feed label information into generator G and discriminator D .

Innovations Better control of generated images; generate tags that are not contained in training set.

Known cons Network model is too simple; only one label for training.

Future work Use more complicated network, add more than one label in training.




Figure 1: Architecture of ss-InfoGAN

1.5 CatGAN

The CatGAN [5] proposes a method for learning a classifier based on unlabeled or partially labeled data.

Abstract

In this paper we present a method for learning a discriminative classifier from unlabeled or partially labeled data. Our approach is based on an objective function that trades-off mutual information between observed examples and their predicted categorical class distribution, against robustness of the classifier to an adversarial generative model. The resulting algorithm can either be interpreted as a natural generalization of the generative adversarial networks (GAN) framework or as an extension of the regularized information maximization (RIM) framework to robust classification against an optimal adversary. We empirically evaluate our method – which we dub categorical generative adversarial networks (or CatGAN) – on synthetic data as well as on challenging image classification tasks, demonstrating the robustness of the learned classifiers. We further qualitatively assess the fidelity of samples generated by the adversarial generator that is learned alongside the discriminative classifier, and identify links between the CatGAN objective and discriminative clustering algorithms (such as RIM).

Problems Existing methods for semi-supervised or unsupervised learning (e.g., Gaussian mixture model, K-means, Maximum margin clustering (MMC), RIM, Boltzman machines, autoencoders) overfit when sample distribution is far from real distribution.

Methods Combining both the generative and the discriminative perspective, learn discriminative neural network classifiers D that maximize mutual information (MI) between the inputs x and the labels y (as predicted through the conditional distribution $p(y|x, D)$) for a number of K unknown categories.

Contributions Introduce MI between observed data x and predicted label y in loss function; provide robustness for RIM.

Future work Introduce Laplacian pyramids.

1.6 BiGAN

The BiGAN [6] learns a generation network and an inference network using an adversarial process.

Abstract

The ability of the Generative Adversarial Networks (GANs) framework to learn generative models mapping from simple latent distributions to arbitrarily complex data distributions has been demonstrated empirically, with compelling results showing that the latent space of such generators captures semantic variation in the data distribution. Intuitively, models trained to predict these semantic latent representations given data may serve as useful feature representations for auxiliary problems where semantics are relevant. However, in their existing form, GANs have no means of learning the inverse mapping – projecting data back into the latent space. We propose Bidirectional Generative Adversarial Networks (BiGANs) as a means of learning this inverse mapping, and demonstrate that the resulting learned feature representation is useful for auxiliary supervised discrimination tasks, competitive with contemporary approaches to unsupervised and self-supervised feature learning.

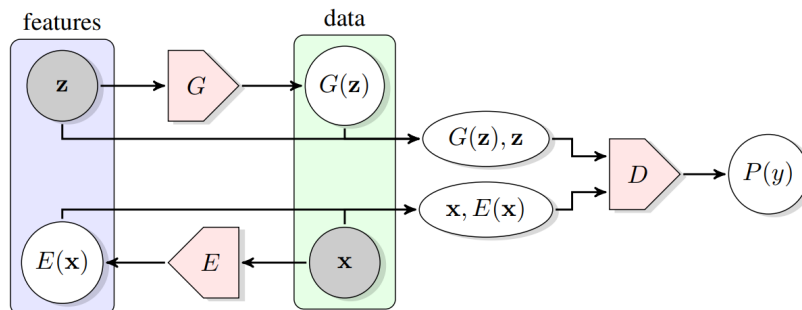


Figure 2: The structure of Bidirectional Generative Adversarial Networks (BiGAN).

Problems Generators of GAN map noise to sample, but unable to do the inverse.

Contributions Not only generating, but also encoding; the learning of the encoder is an auxiliary for the main task.

1.7 ACGAN

The ACGAN [7] improves the performance of generated samples of GANs.

Abstract

In this paper we introduce new methods for the improved training of generative adversarial networks (GANs) for image synthesis. We construct a variant of GANs employing label conditioning that results in 128×128 resolution image samples exhibiting global coherence. We expand on previous work for image quality assessment to provide two new analyses for assessing the discriminability and diversity of samples from class-conditional image synthesis models. These analyses demonstrate that high resolution samples provide class information not present in low resolution samples. Across 1000 ImageNet classes, 128×128 samples are more than twice as discriminable as artificially resized 32×32 samples. In addition, 84.7% of the classes have samples exhibiting diversity comparable to real ImageNet data.

Problems Transform a 32×32 picture to 128×128 will result in blur on picture. Models is not very good if GANs generate only one samples for each class.

Methods Use inception accuracy and MS-SSIM (Multi-Scale Structural Similarity for IMage quality assessment) to assess generated samples.

Innovations Measurement of synthetic images and crash of GAN.

1.8 CVAE-GAN

The CVAE-GAN [8] is a general learning framework that combines VAE and GAN.

Abstract

We present variational generative adversarial networks, a general learning framework that combines a variational auto-encoder with a generative adversarial network, for synthesizing images in fine-grained categories, such as faces of a specific person or objects in a category. Our approach models an image as a composition of label and latent attributes in a probabilistic model. By varying the fine-grained category label fed into the resulting generative model, we can generate images in a specific category with randomly drawn values on a latent attribute vector. Our approach has two novel aspects. First, we adopt a cross entropy loss for the discriminative and classifier network, but a mean discrepancy objective for the generative network. This kind of asymmetric loss function makes the GAN training more stable. Second, we adopt an encoder network to learn the relationship between the latent space and the real image space, and use pairwise feature matching to keep the structure of generated images. We experiment with natural images of faces, flowers, and birds, and demonstrate that the proposed models are capable of generating realistic and diverse samples with fine-grained category labels. We further show that our models can be applied to other tasks, such as image inpainting, super-resolution, and data augmentation for training better face recognition models.

Problems It is difficult for generative model to capture the underlying data distribution since a collection of image samples may lie on a very complex manifold. Also, what is we want to generate images of fine-grained object categories?

Methods Combine CVAE and GAN.

Innovations

- Change generator loss to average divergence, thus stabilizing the training.
- Use encoder thus generating better images.
- Contributes to image inpainting, images rotate and data augmentation.

2 Other works

LiKun Cai Derive a new GAN objective using Renyi Divergence.

Kai Liang Deep learning (DL) theory: why that simple SGD (Stochastic Gradient Descent) do work in DL? From three perspective: 1) Generalization; 2) Optimization; 3) Approximation.

3 From Prof. Geng

These may be Prof. Geng's research interests.

1. Convergence and reachability of multi-user channels.
2. Random code may work powered by today's computation ability, is there any mathematical principle hidden in it?
3. Can multi-letter random code degenerate to single-letter case, is there any structure in it?
4. Assessment of outputs of generator in GAN.
5. Like data which can be compressed, is it possible for network structure?

References

- [1] X. Chen, Y. Duan, R. Houthoofd, J. Schulman, I. Sutskever, and P. Abbeel, "Infogan: Interpretable representation learning by information maximizing generative adversarial nets," in Advances in neural information processing systems, pp. 2172–2180, 2016.
- [2] A. Spurr, E. Aksan, and O. Hilliges, "Guiding infogan with semi-supervision," in Joint European Conference on Machine Learning and Knowledge Discovery in Databases, pp. 119–134, Springer, 2017.
- [3] A. Odena, "Semi-supervised learning with generative adversarial networks," arXiv preprint arXiv:1606.01583, 2016.
- [4] M. Mirza and S. Osindero, "Conditional generative adversarial nets," arXiv preprint arXiv:1411.1784, 2014.
- [5] J. T. Springenberg, "Unsupervised and semi-supervised learning with categorical generative adversarial networks," arXiv preprint arXiv:1511.06390, 2015.
- [6] J. Donahue, P. Krähenbühl, and T. Darrell, "Adversarial feature learning," arXiv preprint arXiv:1605.09782, 2016.
- [7] A. Odena, C. Olah, and J. Shlens, "Conditional image synthesis with auxiliary classifier gans," in Proceedings of the 34th International Conference on Machine Learning-Volume 70, pp. 2642–2651, JMLR.org, 2017.
- [8] J. Bao, D. Chen, F. Wen, H. Li, and G. Hua, "Cvae-gan: fine-grained image generation through asymmetric training," in Proceedings of the IEEE International Conference on Computer Vision, pp. 2745–2754, 2017.